



[Link to Articles](#)

Introduction to Matryoshka Embedding Models

 Upvote **73**

Updated February 23, 2024

[View on GitHub](#)
 [tomaarsen](#)
Tom Aarsen
 [Xenova](#)
Joshua
 [osanseviero](#)
Omar Sanseviero

In this blogpost, we will introduce you to the concept of Matryoshka Embeddings and explain why they are useful. We will discuss how these models are theoretically trained and how you can train them using Sentence Transformers.

Additionally, we will provide practical guidance on how to use Matryoshka Embedding Models and share a comparison between a Matryoshka embedding model and a regular embedding model. Finally, we invite you to check out our interactive demo that showcases the power of these models.

Table of Contents

[Understanding Embeddings](#)

 [Matryoshka Embeddings](#)

 [Matryoshka Dolls](#)

[Why would you use !\[\]\(4186b6ce3a1c83eabb297c1bfd00309c_img.jpg\) Matryoshka Embedding models?](#)

[How are !\[\]\(206536f97fdb267876a3a10ea42b0254_img.jpg\) Matryoshka Embedding models trained?](#)

- [Theoretically](#)

- [In Sentence Transformers](#)

How do I use 🧸 [Matryoshka Embedding models](#)?

- [Theoretically](#)
- [In Sentence Transformers](#)

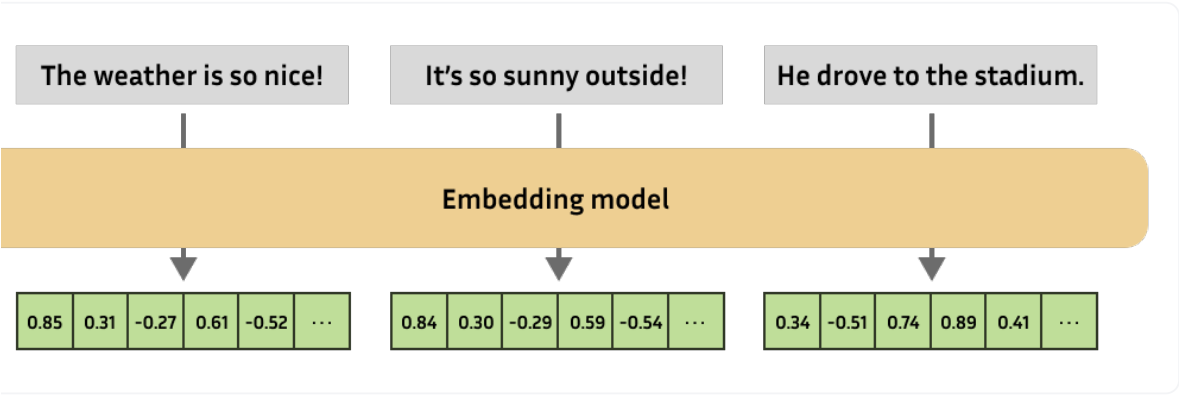
[Results](#)

[Demo](#)

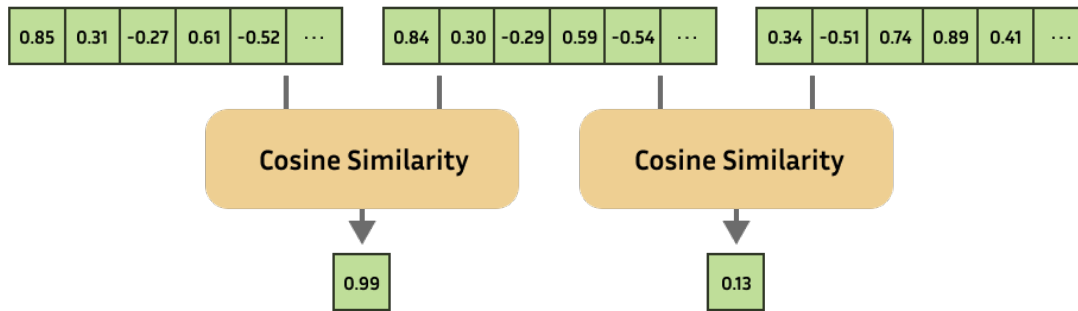
[References](#)

Understanding Embeddings

Embeddings are one of the most versatile tools in natural language processing, enabling practitioners to solve a large variety of tasks. In essence, an embedding is a numerical representation of a more complex object, like text, images, audio, etc.



An embedding model will always produce embeddings of the same fixed size. You can compute the similarity of complex objects by computing the similarity of the respective embeddings!

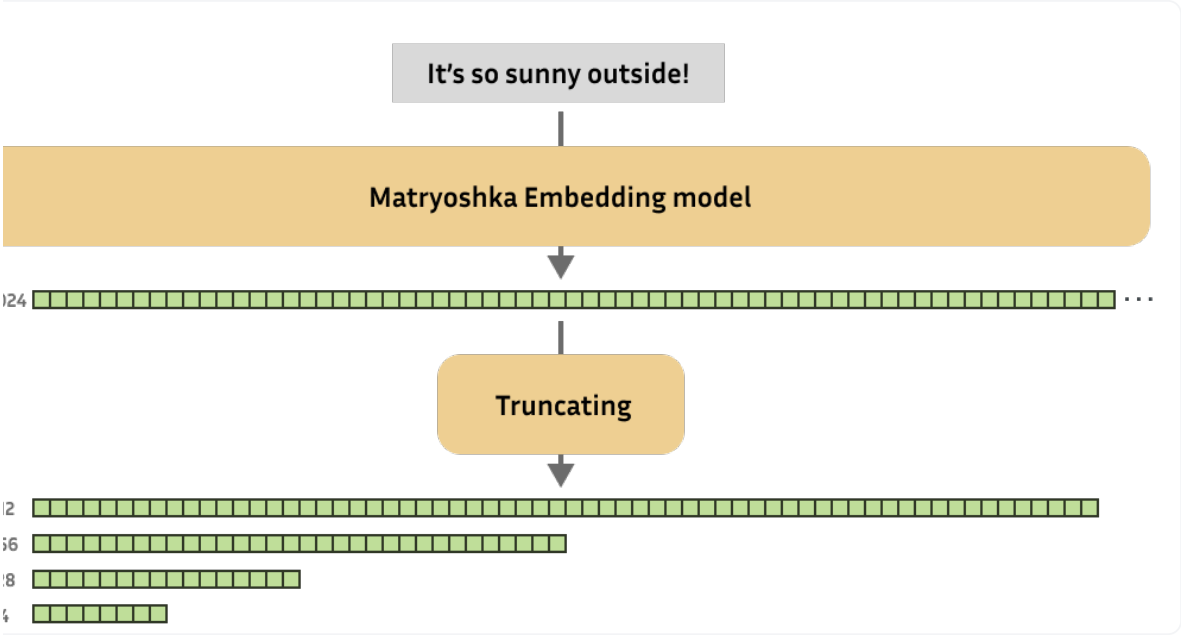


has an enormous amount of use cases, and serves as the backbone for recommendation systems, retrieval, one-shot or few-shot learning, outlier detection, similarity search, paraphrase detection, clustering, classification, and much more!

Matryoshka Embeddings

As search progressed, new state-of-the-art (text) embedding models started producing embeddings with increasingly higher output dimensions, i.e., every input text is represented using more values. Although this improves performance, it comes at the cost of efficiency of downstream tasks such as search or classification.


Recently, [Kusupati et al. \(2022\)](#) were inspired to create embedding models whose embeddings could reasonably be shrunk without suffering too much on performance.



Matryoshka embedding models are trained such that these small truncated embeddings would still be useful. In short, Matryoshka embedding models can produce embeddings of various dimensions.

Matryoshka Dolls

Those unfamiliar, "Matryoshka dolls", also known as "Russian nesting dolls", are a set of wooden dolls of decreasing size that are placed inside one another. In a similar way, Matryoshka embedding models aim to store more important information in earlier dimensions, and less important information in later dimensions. This characteristic of Matryoshka embedding models allows us to truncate the original (large) embedding produced by the model, while still retaining enough of the information to perform well on downstream tasks.



1. Compute Matryoshka Embedding

Why would you use 🧸 Matryoshka Embedding models?

Variable-size embedding models can be quite valuable to practitioners, for example:

Shortlisting and reranking: Rather than performing your downstream task (e.g., nearest neighbor search) on the full embeddings, you can shrink the embeddings to a

smaller size and very efficiently "shortlist" your embeddings. Afterwards, you can process the remaining embeddings using their full dimensionality.

Trade-offs: Matryoshka models will allow you to scale your embedding solutions to your desired storage cost, processing speed, and performance.

How are 🧸 Matryoshka Embedding models trained?

Theoretically

Matryoshka Representation Learning (MRL) approach can be adopted for almost all embedding model training frameworks. Normally, a training step for an embedding model involves producing embeddings for your training batch (of texts, for example) and then applying some loss function to create a loss value that represents the quality of the produced embeddings. The optimizer will adjust the model weights throughout training to reduce the loss value.

For Matryoshka Embedding models, a training step also involves producing embeddings for your training batch, but then you use some loss function to determine not just the quality of your full-size embeddings, but also the quality of your embeddings at various different dimensionalities. For example, output dimensionalities are 768, 512, 256, 128, 64. The loss values for each dimensionality are added together, resulting in a final loss value. The optimizer will then try and adjust the model weights to lower this loss value.

In practice, this incentivizes the model to frontload the most important information at the beginning of an embedding, such that it will be retained if the embedding is truncated.

Sentence Transformers

Sentence Transformers is a commonly used framework to train embedding models, and it natively implements support for Matryoshka models. Training a Matryoshka embedding model using Sentence Transformers is quite elementary: rather than applying some loss function on only the full-size embeddings, we also apply that same loss function on

ated portions of the embeddings.

example, if a model has an original embedding dimension of 768, it can now be trained with losses of 768, 512, 256, 128 and 64. Each of these losses will be added together, optionally with a weight:

```
from sentence_transformers import SentenceTransformer
from sentence_transformers.losses import CoSENTLoss, MatryoshkaLoss

model = SentenceTransformer("microsoft/mpnet-base")

base_loss = CoSENTLoss(model=model)
loss = MatryoshkaLoss(
    model=model,
    base_loss=base_loss,
    matryoshka_dims=[768, 512, 256, 128, 64],
    matryoshka_weight=[1, 1, 1, 1, 1],

model.fit(
    train_objectives=[(train_dataset, loss)],
    ...,
```

Using with MatryoshkaLoss does not incur a notable overhead in training time.

References:

[MatryoshkaLoss](#)

[CoSENTLoss](#)

[SentenceTransformer](#)

[SentenceTransformer.fit](#)

[Matryoshka Embeddings - Training](#)

The following complete scripts as examples of how to apply the MatryoshkaLoss in

ice:

matryoshka_nli.py: This example uses the `MultipleNegativesRankingLoss` with `MatryoshkaLoss` to train a strong embedding model using Natural Language Inference (NLI) data. It is an adaptation of the [NLI](#) documentation.

matryoshka_nli_reduced_dim.py: This example uses the `MultipleNegativesRankingLoss` with `MatryoshkaLoss` to train a strong embedding model with a small maximum output dimension of 256. It trains using Natural Language Inference (NLI) data, and is an adaptation of the [NLI](#) documentation.

matryoshka_sts.py: This example uses the `CoSENTLoss` with `MatryoshkaLoss` to train an embedding model on the training set of the STSBenchmark dataset. It is an adaptation of the [STS](#) documentation.

How do I use 🧸 Matryoshka Embedding models?

Theoretically

In practice, getting embeddings from a Matryoshka embedding model works the same way as with a normal embedding model. The only difference is that, after receiving the embeddings, we can optionally truncate them to a smaller dimensionality. Do note that if embeddings were normalized, then after truncating they will no longer be, so you may need to re-normalize.

After truncating, you can either directly apply them for your use cases, or store them such that they can be used later. After all, smaller embeddings in your vector database should result in considerable speedups!

Keep in mind that although processing smaller embeddings for downstream tasks (e.g., retrieval, clustering, etc.) will be faster, getting the smaller embeddings from the model is as fast as getting the larger ones.

Sentence Transformers

ntence Transformers, you can load a Matryoshka Embedding model just like any other
el, but you can specify the desired embedding size using the `truncate_dim` argument.
that, you can perform inference using the `SentenceTransformers.encode` function, and
mbeddings will be automatically truncated to the specified size.

try to use a model that I trained using `matryoshka_nli.py` with `microsoft/mpnet-base`:

```
from sentence_transformers import SentenceTransformer
from sentence_transformers.util import cos_sim

matryoshka_dim = 64

model = SentenceTransformer("tomaarsen/mpnet-base-nli-matryoshka", truncate_dim=matryoshka_dim)

embeddings = model.encode(
    [
        "The weather is so nice!",
        "It's so sunny outside!",
        "He drove to the stadium.",
    ]
)

print(int(embeddings.shape))
# => (3, 64)

Similarity of the first sentence to the other two:
similarities = cos_sim(embeddings[0], embeddings[1:])
print(similarities)
# => tensor([[0.8910, 0.1337]])
```

Link to the model: [tomaarsen/mpnet-base-nli-matryoshka](https://huggingface.co/tomaarsen/mpnet-base-nli-matryoshka)

free to experiment with using different values for `matryoshka_dim` and observe how
affects the similarities. You can do so either by running this code locally, on the cloud
as with [Google Colab](#), or by checking out the [demo](#).

ences:

[SentenceTransformer](#)

[SentenceTransformer.encode](#)

[util.cos_sim](#)

[Matryoshka Embeddings - Inference](#)

[click here to see how to use the Nomic v1.5 Matryoshka Model](#)

Results


Now that Matryoshka models have been introduced, let's look at the actual performance we may be able to expect from a Matryoshka embedding model versus a regular embedding model. For this experiment, I have trained two models:

[tomaarsen/mpnet-base-nli-matryoshka](#): Trained by running [matryoshka_nli.py](#) with [microsoft/mpnet-base](#).

[tomaarsen/mpnet-base-nli](#): Trained by running a modified version of [matryoshka_nli.py](#) where the training loss is only `MultipleNegativesRankingLoss` rather than `MatryoshkaLoss` on top of `MultipleNegativesRankingLoss`. I also use [microsoft/mpnet-base](#) as the base model.

Both of these models were trained on the AllNLI dataset, which is a concatenation of the [MultiNLI](#) and [MultiNLI](#) datasets. I have evaluated these models on the [STS Benchmark](#) test set using multiple different embedding dimensions. The results are plotted in the following figure:

STSB test score for various embedding dimensions (via truncation), with and without Matryoshka loss



From the top figure, you can see that the Matryoshka model reaches a higher Spearman correlation than the standard model at all dimensionalities, indicative that the Matryoshka model is superior in this task.

Furthermore, the performance of the Matryoshka model falls off much less quickly than the standard model. This is shown clearly in the second figure, which shows the performance at the embedding dimension relative to the maximum performance. **Even at 64, of the embedding size, the Matryoshka model preserves 98.37% of the performance**, much higher than the 96.46% by the standard model.

These findings are indicative that truncating embeddings by a Matryoshka model could: 1) significantly speed up downstream tasks such as retrieval and 2) significantly save on memory space, all without a notable hit in performance.

mo

is demo, you can dynamically shrink the output dimensions of the [nomic-ai/nomic-embed-text-v1.5](#) Matryoshka embedding model and observe how it affects the retrieval performance. All of the embeddings are computed in the browser using [🧠 Transformers.js](#).

Adaptive Retrieval w/ Matryoshka Embeddings

Powered by Nomic Embed v1.5 and 🧠 Transformers.js

Query

What is a panda?

Text

Once upon a time, in a land far, far away...
A panda is a large black-and-white bear native to China.
This is an example sentence.
Hello world.
Ailuropoda melanoleuca is a bear species endemic to China.
The typical life span of a panda is 20 years in the wild.
A panda's diet consists almost entirely of bamboo.
I love pandas so much!
Bamboo is a fast-growing, woody grass.
My favorite movie is Kung Fu Panda.
I love the color blue.

Compute Embeddings

References

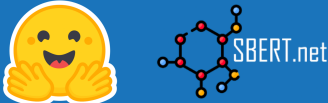
Kusupati, A., Bhatt, G., Rege, A., Wallingford, M., Sinha, A., Ramanujan, V., ... & Farhadi, A. (2022). Matryoshka representation learning. Advances in Neural Information Processing Systems, 35, 30233-30249.
<https://arxiv.org/abs/2205.13147>

Matryoshka Embeddings — Sentence-Transformers documentation. (n.d.).
<https://sbert.net/examples/training/matryoshka/README.html>

UKPLab. (n.d.). GitHub. <https://github.com/UKPLab/sentence-transformers>


Unboxing Nomic Embed v1.5: Resizable Production Embeddings with Matryoshka Representation Learning. (n.d.). <https://blog.nomic.ai/posts/nomic-embed-matryoshka>

More Articles from our Blog



Train 400x faster Static Embedding Models with Sentence Transformers

By tomaarsen January... • △ 126



vdr-2b-multi-v1

Visual Document Retrieval Goes Multilingual

By marco January 9, 2... • △ 64

Community

Preview	
discussing this article	